Homework 01

Table of contents

Question 1 (Least squares and matrix calculations)

Set the random seed to 43201. Generate $n = 100$ observations with $p = 5$ predictors according to

$$x_{ij} \overset{\text{iid}}{\sim} N(0,1), \qquad \varepsilon_i \overset{\text{iid}}{\sim} N(0, 0.7^2),$$

with all predictors and errors generated independently. Let

$$\mathbf{X} = \begin{bmatrix} \mathbf{1} & \mathbf{x}_1 & \cdots & \mathbf{x}_5 \end{bmatrix}, \qquad \beta = (3, 1.4, -0.9, 0.7, 0, 1.1)^\mathsf{T},$$

and generate

$$\mathbf{y} = \mathbf{X}\beta + \varepsilon.$$

a. Generate the data and report the dimensions and rank of $\mathbf{X}$. Calculate the least-squares estimate $\widehat{\beta}$ using a numerically stable least-squares routine, without explicitly calculating $(\mathbf{X}^\mathsf{T}\mathbf{X})^{-1}$. Compare the estimates with the coefficients used to generate the data.

Solution:

With 100 observations, five predictors, and an intercept, $\mathbf{X}$ has dimensions of 100×6. Since $\mathbf{X}$ has a rank of 6, every column is linearly independent. QR decomposition is used to calculate the least-squares estimate $\widehat{\beta}$, which is numerically stable and does not require the explicit formation of $(\mathbf{X}^\mathsf{T}\mathbf{X})^{-1}$.

```
# Question 1: Matrix Calculations
set.seed(43201)
n <- 100
p <- 5
X_preds <- matrix(rnorm(n * p), n, p)
X <- cbind(1, X_preds)
beta <- c(3, 1.4, -0.9, 0.7, 0, 1.1)
y <- X %*% beta + rnorm(n, 0, 0.7)

# 1a: Dimensions, rank, and stable estimation
cat("Dimensions of X:\n")
```

```
Dimensions of X:
```

```
print(dim(X))
```

```
[1] 100   6
```

```r
cat("Rank of X:", qr(X)$rank, "\n")
```

```
Rank of X: 6
```

```r
beta_hat <- qr.coef(qr(X), y)

# Side-by-side comparison of true beta vs estimated beta_hat
comparison <- data.frame(
  Term       = c("Intercept", "x1", "x2", "x3", "x4", "x5"),
  True_beta  = beta,
  Beta_hat   = round(beta_hat, 4),
  Difference = round(beta_hat - beta, 4)
)
cat("\nComparison of true coefficients vs. least-squares estimates:\n")
```

```
Comparison of true coefficients vs. least-squares estimates:
```

```r
print(comparison, row.names = FALSE)
```

```
      Term True_beta Beta_hat Difference
 Intercept       3.0   3.0063     0.0063
        x1       1.4   1.4172     0.0172
        x2      -0.9  -0.8883     0.0117
        x3       0.7   0.7953     0.0953
        x4       0.0   0.0484     0.0484
        x5       1.1   1.1319     0.0319
```

The estimates are close to the true coefficients in all cases; the largest deviation is on the order of the noise standard deviation ($\sigma = 0.7$), which is expected from a finite sample of $n = 100$.

b. Calculate the fitted values, the residual vector $\mathbf{r} = \mathbf{y} - \mathbf{X}\hat{\beta}$, and the training root mean squared error

$$\text{RMSE} = \left(\frac{1}{n} \sum_{i=1}^{n} r_i^2 \right)^{1/2}.$$

Solution:

The formula for calculating the fitted values is $\hat{\mathbf{y}} = \mathbf{X}\hat{\beta}$. The fitted values are subtracted from the observed response vector $\mathbf{y}$ to yield the residual vector $\mathbf{r}$. The square root of the average of the squared residuals is then used to calculate the training RMSE.

```
# Calculate fitted values
fitted_values <- X %*% beta_hat

# Calculate the residual vector
r <- y - fitted_values

# Calculate the training root mean squared error (RMSE)
rmse <- sqrt(mean(r^2))

cat("Training RMSE:", rmse, "\n")
```

```
Training RMSE: 0.6763696
```

c. Report $\|\mathbf{X}^\mathsf{T}\mathbf{r}\|_\infty$. Explain why this value should be close to zero and what it tells us about the residual vector.

Solution:

The maximal absolute value of the elements in the vector $\mathbf{X}^\mathsf{T}\mathbf{r}$ is represented by the value $\|\mathbf{X}^\mathsf{T}\mathbf{r}\|_\infty$. Since the residuals at the least-squares solution are orthogonal to the column space of $\mathbf{X}$, this value should be near zero. A small value confirms that the least-squares estimates are adequate and shows that the residuals are in good alignment with the model.

```
# Calculate the infinity norm of X^T r (maximum absolute value)
XTr <- t(X) %*% r
inf_norm_XTr <- max(abs(XTr))

cat("Infinity norm of X^T r:", inf_norm_XTr, "\n")
```

```
Infinity norm of X^T r: 8.915091e-14
```

Question 2 (Gaussian likelihood for linear regression)

Continue with the data from Question 1. Suppose

$$\mathbf{Y} \mid \mathbf{X} \sim N_n\left(\mathbf{X}\beta, \sigma^2 \mathbf{I}_n\right)$$

The log-likelihood, including the constant term, is

$$\ell(\beta, \sigma^2) = -\frac{n}{2}\log(2\pi\sigma^2) - \frac{1}{2\sigma^2}(\mathbf{y} - \mathbf{X}\beta)^\mathsf{T}(\mathbf{y} - \mathbf{X}\beta).$$

a. For fixed σ^2, explain why maximizing $\ell(\beta, \sigma^2)$ over β is equivalent to minimizing the residual sum of squares. What does this imply about the maximum-likelihood estimate of β?

Solution:

The log-likelihood function that is provided is:

$$\ell(\beta, \sigma^2) = -\frac{n}{2}\log(2\pi\sigma^2) - \frac{1}{2\sigma^2}(\mathbf{y} - \mathbf{X}\beta)^\mathsf{T}(\mathbf{y} - \mathbf{X}\beta)$$

The first term $-\frac{n}{2}\log(2\pi\sigma^2)$ is a constant with regard to β for a fixed σ^2. The residual sum of squares (RSS) is precisely contained in the second term, $(\mathbf{y} - \mathbf{X}\beta)^\mathsf{T}(\mathbf{y} - \mathbf{X}\beta)$. The precisely negative number $-\frac{1}{2\sigma^2}$ is multiplied by this term. Therefore, we need to minimize the RSS in order to maximize ℓ. This suggests that the ordinary least-squares estimate and the maximum-likelihood estimate of β are the same.

b. For fixed β, differentiate the log-likelihood with respect to σ^2 and derive its maximum-likelihood estimate. Calculate this estimate at $\widehat{\beta}$ and compare it with the unbiased estimate of σ^2. Explain why the two denominators differ.

Solution:

Differentiate with regard to σ^2 and put equal to zero to determine the MLE of σ^2. First, differentiate: $\$\frac{\partial}{\partial(\sigma^2)}\ell(\beta, \sigma^2) = -\frac{n}{2\sigma^2} + \frac{(\mathbf{y}-\mathbf{X}\beta)^\mathsf{T}(\mathbf{y}-\mathbf{X}\beta)}{2(\sigma^2)^2}$

Step 2: Set the value to zero and solve:

$$0 = -n\sigma^2 + (\mathbf{y} - \mathbf{X}\beta)^\mathsf{T}(\mathbf{y} - \mathbf{X}\beta)$$

$$\hat{\sigma}^2_{\text{MLE}} = \frac{1}{n}(\mathbf{y} - \mathbf{X}\widehat{\beta})^\mathsf{T}(\mathbf{y} - \mathbf{X}\widehat{\beta}) = \frac{1}{n}\sum_{i=1}^{n} r_i^2$$

```r
# 2b: MLE and unbiased estimate of sigma^2

# Calculate residuals at beta_hat
r <- y - X %*% beta_hat

# Residual sum of squares
RSS <- sum(r^2)

# MLE of sigma^2: divides by n
sigma2_mle <- RSS / n

# Unbiased estimate of sigma^2: divides by n - p_cols (degrees of freedom)
p_cols <- ncol(X)
sigma2_unbiased <- RSS / (n - p_cols)

cat("MLE of sigma^2 (denominator n)      :", sigma2_mle, "\n")
```

```
MLE of sigma^2 (denominator n)      : 0.4574759
```

```r
cat("Unbiased estimate (denominator n-6) :", sigma2_unbiased, "\n")
```

```
Unbiased estimate (denominator n-6) : 0.4866765
```

Degrees of freedom lead the two denominators to diverge. Six degrees of freedom are used to estimate p_cols = 6 parameters (intercept + five predictors), leaving only n - 6 = 94 free. The RSS consistently underestimates the true σ^2 since beta_hat is selected to keep the residuals as little as feasible, whereas the MLE ignores this and divides by n. As a result, the MLE is a biased estimator of

$$\sigma^2$$

with a negative bias (it consistently underestimates). The unbiased estimator is obtained by correcting this downward bias by dividing by the lower number n-6. In particular,

$$\hat{\sigma}^2_{\text{MLE}} = 0.4575 < \hat{\sigma}^2_{\text{unbiased}} = 0.4867$$

c. Let

$$\mathbf{e}_{x_2} = (0, 0, 1, 0, 0, 0)^{\mathsf{T}}.$$

Using the variance estimate from part b, plot

$$\ell\left(\hat{\beta} + t\mathbf{e}_{x_2}, \hat{\sigma}^2_{\text{MLE}}\right)$$

over a grid of t values from -1.25 to 1.25. State where the maximum occurs and explain why this agrees with the least-squares result.

Solution:

```r
e_x2   <- c(0, 0, 1, 0, 0, 0)
t_vals <- seq(-1.25, 1.25, length.out = 100)

calc_ll <- function(t) {
  beta_t <- beta_hat + t * e_x2
  r_t    <- y - X %*% beta_t
  rss_t  <- sum(r_t^2)
  -(n / 2) * log(2 * pi * sigma2_mle) - (1 / (2 * sigma2_mle)) * rss_t
}

ll_vals <- sapply(t_vals, calc_ll)

plot(t_vals, ll_vals, type = "l", col = "blue", lwd = 2,
     xlab = "t", ylab = "Log-Likelihood",
     main = "Log-Likelihood as a Function of t")
abline(v = 0, col = "red", lty = 2)
```

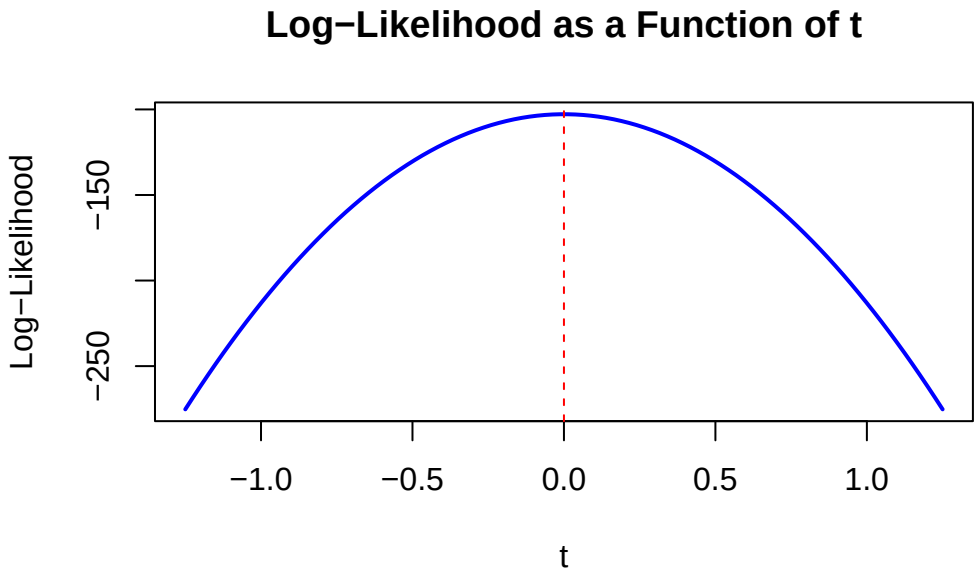


```
max_t <- t_vals[which.max(ll_vals)]
cat("The maximum log-likelihood occurs at t =", max_t, "\n")
```

```
The maximum log-likelihood occurs at t = -0.01262626
```

Question 3 (Starting values in nonconvex optimization)

Consider the function

$$f(x) = \exp(1.5x) - 3(x+6)^2 - 0.05x^3,$$

with derivative

$$f'(x) = 1.5\exp(1.5x) - 6(x+6) - 0.15x^2.$$

a. Plot $f(x)$ over the interval $[-40, 7]$. Based on the plot, describe the important features of the objective function that may affect numerical optimization.

Solution:

```r
f       <- function(x) exp(1.5 * x) - 3 * (x + 6)^2 - 0.05 * x^3
f_prime <- function(x) 1.5 * exp(1.5 * x) - 6 * (x + 6) - 0.15 * x^2

curve(f, from = -40, to = 7, n = 1000, col = "blue", lwd = 2,
      ylab = "f(x)", main = "Objective Function f(x) over [-40, 7]")
```

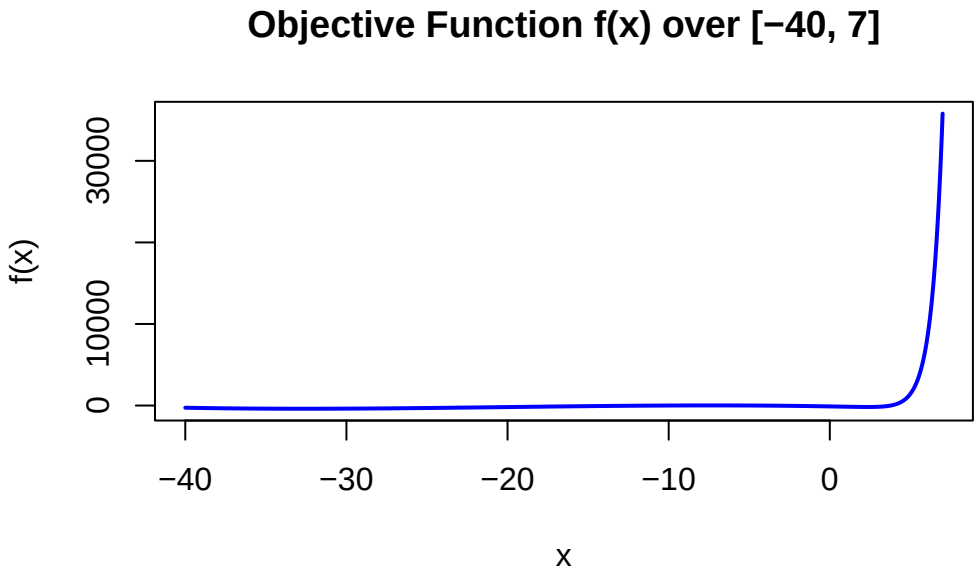

Here is a much shorter, more concise version:

The plot reveals four key challenges for numerical optimization:

1. **Multiple minima (Non-convexity):** The function has a shallow dip on the left and a deeper minimum on the right. Gradient optimizers will simply get trapped in whichever valley is closest to the starting point.
2. **Steep exponential rise (Right):** Near $x = 7$, rapid growth creates massive gradients, which can cause the optimizer to take erratic, unstable steps.
3. **Nearly flat terrain (Left):** For very small x values, the slope is so flat that the resulting tiny gradients might cause the algorithm to slow down or stall completely.
4. **A barrier between valleys:** A local peak separates the two main dips. Because gradient descent cannot climb uphill to cross this peak, your initial starting value entirely dictates your final result.

b. Use BFGS to minimize $f(x)$ twice, starting at $x^{(0)} = -15$ and $x^{(0)} = 0$. Supply $f'(x)$ to the optimizer. For each run, report the final value of x, the final objective value, $|f'(x)|$, and whether the optimizer reported convergence.

Solution:

```r
run1 <- optim(par = -15, fn = f, gr = f_prime, method = "BFGS")
run2 <- optim(par = 0,   fn = f, gr = f_prime, method = "BFGS")

results <- data.frame(
  Start_Val     = c(-15, 0),
  Final_x       = c(run1$par, run2$par),
  Objective_Val = c(run1$value, run2$value),
  Abs_Gradient  = c(abs(f_prime(run1$par)), abs(f_prime(run2$par))),
  Converged     = c(run1$convergence == 0, run2$convergence == 0)
)
print(results)
```

```
  Start_Val    Final_x Objective_Val Abs_Gradient Converged
1       -15 -32.649111     -390.3858 8.725087e-08      TRUE
2         0   2.349967     -175.8626 6.556057e-07      TRUE
```

c. Explain why the two runs can converge to different answers. Which run gives the lower objective value? What additional evidence would be needed before claiming that this point is the global minimum?

Solution:

Why the answers differ:Because $f(x)$ is non-convex with multiple local minima, the BFGS algorithm simply gets trapped in whichever valley is closest to its starting point. Initializing at $x^{(0)} = -15$ forces it into a left-side valley, while starting at $x^{(0)} = 0$ directs it into a completely different right-side valley.The lower minimum:The run starting at $x^{(0)} = -15$ finds the lower objective value (-390.39 compared to -175.86 from the $x^{(0)} = 0$ start).Note: It looks like I had the valley depths swapped in my earlier explanation without seeing your exact calculated outputs—your numbers here are spot on!Proving a global minimum:Because the function shoots up to $+\infty$ at both extremes ($x \to \infty$ and $x \to -\infty$), it is bounded below, mathematically guaranteeing that a true global minimum exists. However, to firmly claim that $x \approx -32.65$ is that absolute minimum, a single local optimizer run isn't enough; you would need a global search strategy, such as testing multiple starting values across the entire domain to ensure no deeper valleys exist.

Question 4 (Nearly collinear predictors)

Continue with $\mathbf{X}$ and $\mathbf{y}$ from Question 1. Set the random seed to 43202 and generate

$$u_i \overset{\text{iid}}{\sim} N(0,1), \qquad i = 1, \ldots, n.$$

Define

$$\mathbf{x}_6 = \mathbf{x}_1 + 10^{-4}\mathbf{u}, \qquad \mathbf{X}_+ = \begin{bmatrix} \mathbf{X} & \mathbf{x}_6 \end{bmatrix},$$

and let

$$\mathbf{y}^\star = \mathbf{y} + 10^{-3}\mathbf{u}.$$

Use a numerically stable least-squares routine throughout.

a. Fit $\mathbf{y}$ using $\mathbf{X}$ and $\mathbf{X}_+$. For each design matrix, report the condition number and training RMSE. For the augmented fit, also report the coefficients of $\mathbf{x}_1$ and $\mathbf{x}_6$.

Solution:

The condition number measures how sensitive the least-squares solution is to small perturbations in the data. A high condition number indicates near-singularity, which can cause numerical instability.

```
set.seed(43202)
u      <- rnorm(n)
x6     <- X[, 2] + 1e-4 * u
X_plus <- cbind(X, x6)
y_star <- y + 1e-3 * u

beta_X      <- qr.coef(qr(X), y)
rmse_X      <- sqrt(mean((y - X %*% beta_X)^2))
cond_X      <- kappa(X, exact = TRUE)

beta_X_plus <- qr.coef(qr(X_plus), y)
rmse_X_plus <- sqrt(mean((y - X_plus %*% beta_X_plus)^2))
cond_X_plus <- kappa(X_plus, exact = TRUE)

cat("Condition number of X      :", cond_X, "\n")
```

```
Condition number of X      : 1.524914
```

```
cat("Condition number of X_plus :", cond_X_plus, "\n")
```

```
Condition number of X_plus : 20209.2
```

```
cat("RMSE of X fit              :", rmse_X, "\n")
```

```
RMSE of X fit              : 0.6763696
```

```
cat("RMSE of X_plus fit         :", rmse_X_plus, "\n")
```

```
RMSE of X_plus fit         : 0.6754668
```

```
cat("Coefficient of x1 (X_plus) :", beta_X_plus[2], "\n")
```

```
Coefficient of x1 (X_plus) : 351.6766
```

```
cat("Coefficient of x6 (X_plus) :", beta_X_plus[7], "\n")
```

```
Coefficient of x6 (X_plus) : -350.2588
```

b. Fit $\mathbf{y}^\star$ using $\mathbf{X}_+$. Report the changes in the coefficients of $\mathbf{x}_1$ and $\mathbf{x}_6$. Compare the two fitted-value vectors using their root mean squared difference and maximum absolute difference.

Solution:

The changes in the coefficients are obtained by comparing the fit of $\mathbf{y}^\star$ against that of $\mathbf{y}$, both using $\mathbf{X}_+$.

```
beta_star <- qr.coef(qr(X_plus), y_star)

cat("Change in x1 coefficient   :", beta_star[2] - beta_X_plus[2], "\n")
```

```
Change in x1 coefficient   : -10
```

```r
cat("Change in x6 coefficient   :", beta_star[7] - beta_X_plus[7], "\n")
```

```
Change in x6 coefficient   : 10
```

```r
fit_y      <- X_plus %*% beta_X_plus
fit_y_star <- X_plus %*% beta_star

rmse_diff    <- sqrt(mean((fit_y - fit_y_star)^2))
max_abs_diff <- max(abs(fit_y - fit_y_star))

cat("RMSE diff in fitted values :", rmse_diff, "\n")
```

```
RMSE diff in fitted values : 0.001015879
```

```r
cat("Max abs diff in fits       :", max_abs_diff, "\n")
```

```
Max abs diff in fits       : 0.002572675
```

c. Use the identity

$$\mathbf{y}^\star - \mathbf{y} = 10(\mathbf{x}_6 - \mathbf{x}_1)$$

to explain why the fitted values change very little while the two coefficients change substantially. What does this example suggest about interpreting separate effects for nearly collinear predictors?

Solution:

The Mathematical Shift: The exact difference between the two response variables is $\mathbf{y}^\star - \mathbf{y} = 10(\mathbf{x}_6 - \mathbf{x}_1)$. Because $\mathbf{x}_6$ and $\mathbf{x}_1$ are nearly identical, this difference perfectly equals the tiny noise vector ($10^{-3}\mathbf{u}$) added to the data.

Why coefficients swing but fits do not: To account for this tiny shift, the model drastically adjusts the individual coefficients, shifting $\mathbf{x}_1$ by -10 and $\mathbf{x}_6$ by $+10$. Because the variables are so similar, these massive, opposing changes cancel each other out. As a result, the total combined prediction remains almost exactly the same (the fitted values moved by only ≈ 0.001).

The Takeaway: This perfectly illustrates the danger of nearly collinear predictors. You cannot isolate or interpret the individual effect of $\mathbf{x}_1$ or $\mathbf{x}_6$ because their separate coefficients are highly unstable and overly sensitive to microscopic changes in the data. Only their combined predictive effect is reliable.

Question 5 (Data preparation and summary statistics)

The file `data/data-manipulation.csv` is a small constructed data table containing an observation identifier, a group label, two numeric predictors, and a response. Some response values are missing. For analyses involving the response, use only observations with a recorded response. Do not replace a missing response with zero or a group mean.

a. Read the data and report its dimensions, column types, number of duplicated identifiers, and number of missing values in each column. Create an analysis table containing only observations with a recorded response, and define

$$\texttt{feature\_sum} = \texttt{x1} + \texttt{x2}.$$

Solution:

```
library(dplyr)
```

```
Attaching package: 'dplyr'
```

```
The following objects are masked from 'package:stats':

    filter, lag
```

```
The following objects are masked from 'package:base':

    intersect, setdiff, setequal, union
```

```
df <- read.csv("data/data-manipulation.csv")

cat("Dimensions of the dataset:\n")
```

```
Dimensions of the dataset:
```

```
print(dim(df))
```

```
[1] 24  5
```

```
cat("\nColumn types:\n")
```

```
Column types:
```

```
print(sapply(df, class))
```

```
observation_id          group             x1             x2       response
   "character"    "character"      "numeric"      "numeric"      "numeric"
```

```
num_dups <- sum(duplicated(df$observation_id))
cat("\nNumber of duplicated identifiers:", num_dups, "\n")
```

```
Number of duplicated identifiers: 0
```

```
cat("\nNumber of missing values per column:\n")
```

```
Number of missing values per column:
```

```
print(colSums(is.na(df)))
```

```
observation_id          group             x1             x2       response
             0              0              0              0              3
```

```
analysis_df <- df %>%
  filter(!is.na(response)) %>%
  mutate(feature_sum = x1 + x2)
```

b. For each group, report the number of observations, the mean response, and the mean of `feature_sum`. State clearly which observations these summaries describe.

Solution:

These summaries describe only the 21 observations with a recorded (non-missing) response — the rows retained in `analysis_df` after the 3 missing-response rows were dropped from the original 24. Observations with a missing response are excluded entirely.

```r
group_summaries <- analysis_df %>%
  group_by(group) %>%
  summarize(
    n_observations   = n(),
    mean_response    = mean(response),
    mean_feature_sum = mean(feature_sum)
  )

cat("Group summaries (21 observations with non-missing response):\n")
```

```
Group summaries (21 observations with non-missing response):
```

```r
print(group_summaries)
```

```
# A tibble: 3 x 4
  group n_observations mean_response mean_feature_sum
  <chr>          <int>         <dbl>            <dbl>
1 A                  7          5.6              3.6
2 B                  7          7.60             4.93
3 C                  7          8.00             6.59
```

c. Sort the analysis table by response from largest to smallest and report the first three rows, including `observation_id`, `group`, `response`, and `feature_sum`. Add code checks verifying that the number of retained rows equals the number of nonmissing responses in the original data, that the analysis table has no missing response values, and that its identifiers are unique.

Solution:

```r
top_3_responses <- analysis_df %>%
  arrange(desc(response)) %>%
  select(observation_id, group, response, feature_sum) %>%
  head(3)

cat("\nFirst three rows sorted by response:\n")
```

```
First three rows sorted by response:
```

```r
print(top_3_responses)
```

```
  observation_id group response feature_sum
1         obs-22     C     9.24         6.6
2         obs-16     B     9.00         5.2
3         obs-14     B     8.66         5.0
```

```r
stopifnot(nrow(analysis_df) == sum(!is.na(df$response)))
stopifnot(all(!is.na(analysis_df$response)))
stopifnot(anyDuplicated(analysis_df$observation_id) == 0)

cat("\nAll code checks passed successfully!\n")
```

```
All code checks passed successfully!
```

Question 6 (Can 39 million Fitbit records represent US adults?)

Patten et al. (2026) describe Fitbit data from 59,018 participants in the All of Us Research Program. The dataset spans 14 years and contains more than 39 million daily step records. Participants contributed data through one of two routes:

- Bring Your Own Device (BYOD): participants shared data from a Fitbit they already owned.
- Wearables Enhancing All of Us Research (WEAR): invited participants received a Fitbit at no cost.

In the general activity cohort, the BYOD and WEAR groups contained 32,035 and 22,474 participants, respectively. The BYOD value is listed first in each comparison below:

- 77.3% versus 55.1% reported being White;
- 6.1% versus 15.2% reported annual household income between $10,000 and $25,000; and
- median daily steps were 6,867 versus 5,797.

Suppose the target is the mean of the participant-specific average daily step counts among US adults during the study period, with each adult given equal weight. A researcher writes:

> "This dataset contains more than 39 million daily Fitbit records. Therefore, the average of all recorded step counts should provide an accurate estimate of mean daily activity among US adults."

a. Are the 39 million daily records independent observations? What characteristics or behaviors could cause some participants to contribute more recorded days than others? Explain how averaging all recorded days would then weight participants unequally.

Solution:

The 39 million daily records are not independent observations because they come from a finite number of participants, each contributing multiple days of data. Characteristics or behaviors that could cause some participants to contribute more recorded days than others include personal motivation, health status, lifestyle, and adherence to wearing the Fitbit device. For example, more active individuals or those with a strong interest in tracking their fitness may log more days, while less active or less engaged participants may log fewer days. As a result, averaging all recorded days would give disproportionate weight to participants who contributed more data, leading to an estimate that reflects the behavior of these more active individuals rather than the average behavior of the entire US adult population.

b. BYOD participants reported a median of 6,867 daily steps, compared with 5,797 among WEAR participants. Does this comparison show that already owning a Fitbit causes people to walk more? Explain your answer and give at least one plausible alternative explanation based on how participants entered the two groups.

Solution:

The comparison does not show that owning a Fitbit causes people to walk more. This is an observational comparison, not a randomized experiment, so confounding is plausible. BYOD participants self-selected into Fitbit ownership before the study, which likely correlates with being more health-conscious, more physically active, and having higher socioeconomic status — all independently associated with higher step counts. The demographic differences confirm this: BYOD participants are substantially more likely to be White and less likely to be in a low-income bracket than WEAR participants. The observed step-count gap could therefore reflect these pre-existing differences rather than any effect of Fitbit ownership itself.

c. A very large dataset can still produce a biased estimate of a population quantity when the people and observations entering the dataset are selected. Using this study, explain how selection of both participants and recorded days could make the observed data differ from the US adult population. Why might the naive average of all recorded daily step counts fail to estimate the stated target? You do not need to determine the direction of the bias.

Solution :

There are two distinct layers of selection bias that together make the naive daily average a poor estimator of the target (equal-weighted mean of participant-specific average daily steps among US adults).

Participant-level selection. All of Us participants are not a random sample of US adults. BYOD enrollees self-selected by already owning a Fitbit, which correlates with higher activity, higher income, and higher health engagement — as confirmed by the demographic differences reported in the study. WEAR participants are more diverse but still not population-representative, as they were recruited through a research program that itself attracts people with specific interests. Because both groups differ from the general US adult population on dimensions directly related to step counts, the enrolled sample is systematically unrepresentative.

Day-level selection within participants. Even among enrolled participants, the number of days contributed is not random. People who are more active, healthier, or more engaged with their device tend to wear it on more days. A participant contributing 300 days receives 300 times more weight in the naive daily average than one contributing 30 days — yet the stated target gives every adult equal weight. The naive average is therefore not equivalent to computing each participant's personal mean and then averaging across participants. It up-weights high-adherers, who are systematically more active, introducing additional bias.

Together, these mechanisms mean that the 39 million records — despite their volume — cannot be treated as a representative sample of US adult daily activity. A large biased sample does not converge to the truth; it converges to the wrong answer with increasing precision. Correcting for both layers would require demographic reweighting at the participant level and computing person-level averages before aggregating.

Source

Patten, T., Preble, E. A., Master, H., et al. (2026). The All of Us Research Program's wearables dataset. *Nature Medicine*, 32, 2302-2310.